

Chaîne de Publication

How the Observatoire works, from a candidate's speech to a line on a public page.

01 The idea

A candidate says something about artificial intelligence. It might be in a fifty page programme, a speech filmed at a tech conference, a post on X, or a vote in parliament. Six months later nobody can find it, and three newspapers have reported three different numbers.

The Observatoire fixes one narrow thing. It watches a fixed list of places where candidates speak, pulls out what they said about AI, and publishes each position with the exact quote, the date and a link. In French and in English.

You read it two ways. **Le fil** is a feed, newest first, for what changed this week. **La grille** is a table, candidates down the side and six policy themes across the top, for comparing them side by side.

The hard part is not collecting the text. It is that an institution's name is on the result. So every rule below exists for one reason, which is that **a reader has to be able to check anything we publish without taking our word for it.**

02 The words we use

Six words that come up on every page below. Two minutes here saves confusion later.

Document

A piece of text we collected and kept.
A programme, an article, the subtitles
of a video.

Lead

Something we noticed but will not
publish from. A headline saying a
candidate spoke about AI. It is a
question for the editor, not an answer.

Position

One thing a candidate said. A quote, a
theme, a date and a link. This is what
appears on the site.

Tier

Whose words they are.

1

the candidate himself,

2

his party,

3

a journalist reporting on him. Not who
hosts the page.

Axis

One of six policy themes: regulation,
sovereignty, work, public services,
surveillance, funding. Every claim
gets exactly one.

Le fil · La grille

The two views of the same claims. A
feed and a comparison table.

03 Sources

What matters is **whose words they are**, not which website they sit on. A speech is the candidate's own words wherever it is published. Only his words and his party's become positions. A journalist writing *about* him becomes a lead instead, which is a question for the editor rather than something we publish. A press article is read once and filed under whichever candidate it names, by looking for the name. We never ask the model who was being quoted.

TIER 1

Own words

→ claims

📄

Programme PDF

his published programme

📄

Page de campagne

a campaign page, revisited each time

🗳️

Scrutins

how he voted, not what he said

🗨️

X timeline

his own account, never a search

🎥

Speech video

a filmed speech, one voice only

TIER 2

Party & staff

→ claims

📡

republicains.fr

the party news feed

📡

lafranceinsoumise

the party news feed

📡

parti-socialiste.fr

the party news feed

🗣️

@lesrepublicains

several voices, so a lead

TIER 3

Press

→ leads only

📡

synthmedia.fr

specialist press, freely readable

📡

next.link

specialist press, freely readable

🔒

Paywalled majors

only the opening is visible, so we quote briefly and link

🗨️ WHERE SOCIAL MEDIA STANDS

X

Reading it costs money, charged per message. Only a candidate's own account, never a search, because a search would sweep in ordinary people.

YouTube

Usable, but several people usually talk in one video, so it becomes a lead rather than a position.

Bluesky, Mastodon

Easy to read, but French candidates are barely there.

TikTok, Instagram

Biggest audience, no lawful way to read them automatically, and short video carries personality rather than policy.

So today no social media produces a claim. That is one deliberate choice and one open blocker.

🚫 TESTED AND RULED OUT

GDELT DOC 2.0

returned nothing at all when we tested it

nosdeputes.fr

the service answers, but with nothing in it

Bluesky · Mastodon

French politics is not there

Europresse / Factiva

the licence forbids automated reading

numerama · usine-digitale

these sites ask not to be read automatically

Academic papers

they do not contain what candidates said

Jina · Firecrawl

they meet the same paywalls we do

vie-publique.fr

it archives *government* speech, and our candidates are not in government

Google News bodies


every link leads to a cookie wall, not the article

TikTok · Instagram

no lawful automated access


04 Infrastructure

Five places. Only one costs money and only one holds a password. Everything else is free.




05 The morning

Time runs left to right, and who does the work runs top to bottom. The gap in the middle is waiting. The work is sent off and collected later rather than sitting idle.



What a site has already published does not change, so this is done once per site and never repeated. A news feed only carries the last few weeks, so without this pass an older speech stays invisible. That is exactly what happened with one candidate's July speech on artificial intelligence.



06 Anatomy of a claim

Everything we publish keeps three voices apart and shows the reader which is which. That separation is the whole argument for trusting it.

» HIS WORDS

« *Pour les Français, je propose un agent IA administratif unique, « Marianne », capable de simplifier toutes leurs démarches.* »

✦ WHAT HE COMMITTED TO

Propose un agent conversationnel unique, « Marianne », pour l'ensemble des démarches administratives.

🔗 WHAT IT MEANS

Guichet unique piloté par IA. Le périmètre des démarches couvertes n'est pas précisé.

📄 THE RECEIPT

theme: public services · his party's site · 7 July 2026 · archived copy kept

» THE CANDIDATE

Exactly as he said it. **Never** edited, corrected or translated. Kept short, because French law allows quoting briefly when you say where it came from.

✦ EXTRACTION

One plain sentence saying what he committed to. It describes and never grades. Words of judgement are rejected automatically.

🔗 METALAB

What the position means in practice. This is where the institute adds something, and it carries a label so no reader mistakes it for his words.

📄 RECORDED AUTOMATICALLY

Where it came from, the date, and a link to an archived copy taken the day we read it. Pages get edited and deleted, so without that copy a quote becomes unprovable.

07 Approaches we rejected

Both were sensible suggestions. Neither survives what this publication has to promise a reader.

Sentiment analysis and EDA

The text is already concrete

Candidates give real figures and named plans. Billions of euros, numbers of civil servants, a scheme with a name. There is nothing to infer. You copy what is written.

Counting has one use, and it stays internal. Seeing which candidate and which theme we have little on tells the editor where to look next. It is never published as a finding.

A score is a judgement

Sentiment produces a number. A number puts candidates in an order. An institution publishing an order of candidates during an election is the one thing everything here exists to avoid.

It cannot be checked

A quote either appears in the source or it does not, and anyone can look. A score has no equivalent test, so a reader would simply have to trust us.

A vision model reading the video

We have not tried it, so we cannot tell you what it would cost. Any figure we gave would be arithmetic rather than experience. None of the reasons below depend on the answer.

It answers a different question

It tells you who is *on screen*. We need to know who is *talking*. French television cuts to the interviewer while the candidate is still speaking, and uses split screens and stock footage. Being on screen usually means you are speaking. Usually is not good enough when the whole product is knowing who said what.

Wrong sense entirely

Telling voices apart is a sound problem, not a picture problem. Tools exist that separate speakers from the audio. That is the right instrument for the job.

It cannot produce the evidence

A quote has to appear word for word in the source. A model looking at pictures describes what it saw, it does not transcribe. So you would still need the subtitles anyway.

It breaks replay

How often you look at the video is a choice. Look at different moments and you get a different answer. Everything here is built so a result can be traced back and repeated.

It saves the reviewer nothing

Someone still confirms before anything is published, and the subtitles already give the exact second, reliably and for nothing.

Where it would genuinely earn its place. Text shown on screen and never spoken, such as a figure that appears only on a slide. And pages of a programme that are charts rather than words. That is a handful of images, not thousands, and it is cheap. Worth doing. Just not for working out who spoke.